

Relational Databases

15

Key Topics



- Tables & Relations
- Primary and Foreign Keys
- Referential Integrity
- Relational Operations
- Data Retrieval Approach

16

Relational Database

A relational database S is a set of relations such that $S = (R_1, R_2, \dots, R_n)$ and $R_i(A_{i1}, A_{i2}, \dots, A_{ik})$.

Each relation instance must satisfy the integrity constraints specified in the relational database schema.

17

Fundamental Structure: Table

Account

acct_no	balance	rate	cust_id
1009845	987.23	2.5	10097

Tables contain data records for a single business domain *entity*

Foreign keys establish links between tables

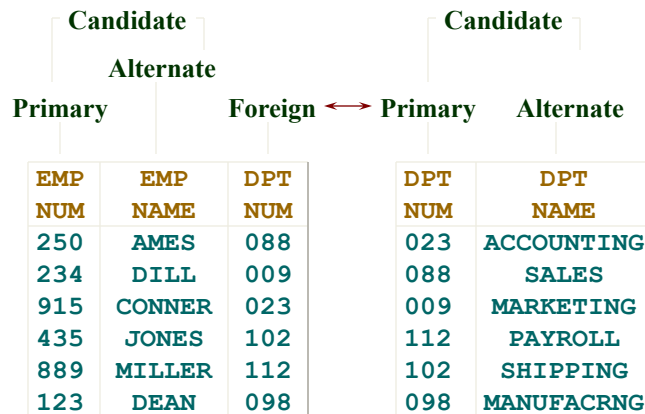
Primary keys uniquely identify rows

Customer

cust_id	cust_name	cust_phone	...
10097	John Smith	(508) 555-1398	

18

Types of Keys



19

Database Integrity

- There are four types of integrity constraints for relational databases:
 - domain integrity
 - entity integrity
 - referential integrity
 - semantic integrity

20

Domain Integrity

- Domain integrity is applied on the attribute and requires that every attribute A_i is a value from its domain D_i .
- More formally, for a relation $R(A_1, A_2, \dots, A_n)$, each instance of the relation satisfies:

$$\{ \langle A_1:V_1, A_2:V_2, \dots, A_n:V_n \rangle \mid V_1 \in D_1, V_2 \in D_2, \dots, V_n \in D_n \}$$

21

Entity Integrity

- Entity integrity defines restrictions on the tuples of a relation which ensures the accuracy and consistency of the database.
- It requires that each tuple in a relation is unique.
- The primary key must be a set of attribute(s) that cannot have duplicate values or be *null*.

22

Primary Key Properties

- Primary keys must be:
 - *unique*
 - *not reusable*
 - *unchangeable*
 - *void of semantic meaning*
 - *invisible to the user*

Defining Primary Keys

- A primary key is a subset of attributes whose values are unique for a relation.
- Often there is no subset of the attributes that satisfies all the above constraints, so a *surrogate* or *artificial* key is created as a new attribute.
- For some relations there is a single attribute that is unique: *natural key*.

Referential Integrity

Account

acct_no	balance	rate	cust_id
1009845	987.23	2.5	10097

Foreign keys establish links between tables; a foreign key must "point" to an existing row with that primary key

Referential integrity (RI) ensures that there are no *dangling references*

Customer

cust_id	cust_name	cust_phone	...
10097	John Smith	(508) 555-1398	

25

Semantic Integrity

- Semantic integrity ensures that business and other logical rules are enforced.
- Examples of semantic integrity rules:
 - Number of projects a project manager manages is greater than 0
 - A project cannot have a start date after its end date
 - A project must have a project manager
 - A project manager cannot be a resource on the project that the project manager manages

26

Relational Table Rules

The following rules govern tables:

1. All rows are formatted identically
2. Rows are unsorted
3. Columns are in no particular order
4. Each row in a table is a unique instance
5. There are no duplicate column (attribute) names
6. Within the table, values represent data
7. Each column (attribute) value is atomic
8. Primary keys should be non-compound and minimal



Understanding Tables and Keys

1. Consider the table *Sections* below.
 - a. What is the likely primary key?
 - b. Is the primary key natural or surrogate?
 - c. What kind of a key is “*instructor*”?

Table: Sections

crn	title	hours	instructor
1009845	DS4100	4.0	10097
1009861	DS4400	4.0	10097
1009893	CS1200	1.0	10461



Creating a Relational Schema

Suppose that you need a relational database for an investment company that must track portfolios, investments (such as stocks, bonds, CDs, etc.), and trades in the portfolios, *e.g.*, buying some stock at some price on some day. Sketch a data model for this with some necessary attributes and keys. Map the model to a database schema, *i.e.*, tables, primary keys, foreign keys.

29



Summary, Review, & Questions...

30